

SwitchSync Markets: A Preregistered Synthetic Stress Test of Switching Synchronization as a Financial-Network Mechanism

SwitchSync Markets Project*

July 2026

Abstract

Dynamic link switching can induce stable synchronized states in sparse multilayer networks of chaotic oscillators [12]. We ask whether this mechanism admits a mathematically faithful and observationally identifiable translation to financial temporal networks, where cross-venue arbitrage channels open and close over time. Rather than fitting market data, we preregistered falsifiable gates and executed them *exactly once* on frozen synthetic systems: a two-layer FitzHugh–Nagumo (FHN) benchmark and a linear stochastic surrogate with paired switching schedules, under a frozen statistical decision contract (exact sign test plus a preregistered effect band). The calibrated demonstration gate passed: at $\sigma_{12} = 1.5$, $N = 40$, fast switching synchronizes the layers and slow switching does not. Every translation-relevant gate, however, failed to establish its claim. The switching-versus-static advantage was directionally consistent (8/8 positive paired differences, $p = 0.0078$) but below the preregistered magnitude band (INCONCLUSIVE/TIE); the strict comparison against the average graph and the best frozen static comparator is not interpretable under the frozen hierarchy, with diagnostically negative differences in all eight seeds; no temporal-order effect was resolved; robustness stages never exceeded their bands; and observational identifiability failed outright (precision and recall ≈ 0.35 against a 0.6 bar), with asynchronous, last-observation-carried-forward sampling collapsing the contraction correlation from 0.948 to 0.010 (an Epps-like degradation [10]). We conclude that the financial translation is not supported for empirical or operational development under the preregistered criteria. This negative result does not refute the mathematics of the seed paper; it constrains what that mechanism can be expected to deliver as a measurable, decision-relevant object in financial data. No real market data were used.

1 Introduction

Eser, Medeiros, Riza and Engel recently showed that switching a sparse set of inter-layer links — rather than keeping them fixed — can stabilize inter-layer synchronization in two-layer networks of chaotic FitzHugh–Nagumo oscillators, with faster switching more effective [11, 12]. The mechanism is attractive to a financial modeler: markets trade the *same* economic asset on multiple venues; the cross-venue “links” (arbitrage channels, shared market makers, funding connections) are sparse, come and go over time, and venue prices are observed to lock together or drift apart episodically [1, 3, 17]. If switching-induced synchronization survived translation, “which channels are open, in what order, and at what rate” would be a structural, measurable determinant of cross-venue price locking, distinct from aggregate connectedness measures [5, 7, 8].

*Repository: `switchsync-markets`. This is a fully synthetic, preregistered study; no real market data were used at any stage. Contact via the repository.

That conditional is precisely what this paper tests, and the honest answer is negative under our preregistered criteria. Three failure modes make naive translations of synchronization results to finance unreliable: (i) fast-switching averaging theory implies that rapidly switched networks behave like their time-averaged counterpart [4, 18, 30], so an apparent “switching effect” may carry no information beyond the average graph; (ii) small ensembles invite effect-size inflation, so magnitude criteria must be fixed before observing data; and (iii) even a real mechanism is useless to empirical finance if it cannot be identified from asynchronously observed prices [10, 19, 29]. We therefore preregistered a gate hierarchy with frozen grids, seeds, estimands and decision rules; froze the code and configuration under content-addressed manifests; executed the suite exactly once under an auditable custody protocol; and committed the sealed outputs byte-for-byte.

Section 7 reports the outcomes (Table 1, Figure 1): the mechanism demonstrably operates in a calibrated synthetic setting (GOB PASS), but no translation-relevant gate established its claim, and identifiability failed. Throughout, we separate *confirmatory* preregistered outcomes, *diagnostic* quantities that the frozen hierarchy declares non-interpretable, *descriptive* observations, and *conjecture*.

2 Mathematical Background

Synchronization and master stability. For diffusively coupled identical oscillators, the stability of the synchronous state is governed by the transverse (conditional) Lyapunov exponents of the variational dynamics [26, 27]. On a static graph the problem factorizes into a master stability function (MSF) evaluated at the Laplacian eigenvalues.

Time-varying coupling. When the coupling network varies in time, the static MSF picture survives exactly only in special cases, e.g. commuting coupling matrices [6]. In the fast-switching regime, averaging results guarantee synchronization whenever the *time-averaged* network synchronizes [4, 30], with finite-rate deviations bounded by blinking-system theory [18]. Consensus theory gives the discrete-time counterpart: joint connectivity over time windows suffices [21, 24, 25, 28]. Empirically, rapidly switched links promote synchronization in a range of models [23, 32]; see [15] for a survey, [20] for the temporal-network formalism, and [22] for rigorous conditions under time-dependent diffusive coupling. *Crucially*, temporal networks *can* outperform any static network of equal resource budget in designed settings [31], so “switching never beats the average” is not a theorem; whether it holds in a given ensemble is an empirical question — the one our G1/G2 gates formalize.

The seed mechanism. The seed paper [12] couples two ring layers of chaotic FHN oscillators through a sparse, randomly re-drawn set of inter-layer links with dwell time T_{swt} , and reports that sufficiently fast switching induces stable inter-layer synchronization at link densities where static sparse coupling fails; its predecessor [11] maps the edges of inter-layer synchronization under time-switched links.

3 Financial-Network Translation and Its Assumptions

The proposed reading maps: layers $\rightarrow$ venues quoting the same asset; inter-layer synchronization $\rightarrow$ cross-venue price locking (conceptually a cointegration-like common component [16]); the sparse switching link set $\rightarrow$ the currently active arbitrage/market-making channels; the switching rate $\rightarrow$ channel-rotation speed. Under this reading, price-discovery leadership that migrates between venues over time [1, 3, 17] is the observable trace of channel switching.

This translation embeds strong assumptions, each of which can independently fail without contradicting the seed mathematics:

1. **Mechanism transfer:** the contraction advantage of switching must survive in dynamics that are not two-layer chaotic FHN rings — here probed with a linear stochastic surrogate (Section 5).
2. **Beyond-average content:** the switching structure must matter *beyond* the time-averaged connectivity; otherwise standard aggregate connectedness measures [7–9] already capture everything, and correlation-based “regime” readings risk the classic interdependence/heteroskedasticity confound [14].
3. **Identifiability:** the active channel set must be recoverable from observed prices, after removing common factors [2, 13] and despite asynchronous observation [10, 19, 29].

We emphasize what the translation, even if successful, would *not* give: a trading signal, a backtest, or any claim about market causality. None is made anywhere in this paper.

4 Hypotheses and Preregistered Gates

The preregistration (v8; hashes in Appendix A) freezes a gate hierarchy. Each gate has a frozen estimand, grid, seed block and decision rule; outcomes are PASS / FAIL / INCONCLUSIVE / EXECUTION_INVALID, with NOT_INTERPRETABLE for sub-gates whose interpretability is conditional.

G0A (exact reproduction; expensive). Reproduce the paper-scale regime ($\sigma_{12} = 0.1$, $N \in \{100, 200, 400\}$). *Not run* in this study (Section 7).

G0B (calibrated demonstration). At a re-calibrated coupling ($\sigma_{12} = 1.5$, $N = 40$), fast switching ($T_{\text{swt}} \leq 10$) must synchronize and slow switching ($T_{\text{swt}} \geq 160$) must not. *This is a qualitative demonstration of the mechanism, never an exact reproduction.*

G0C (minimal MSF). In a minimal two-node transverse-stability computation with switched anti-phase coupling channels, the stability onset must depend on T_{swt} .

G1-weak. In the surrogate, the selected switching arm must beat the equal-instantaneous-density static-sparse comparator on paired evaluation seeds.

G1-strict. The arm must beat *both* the average-graph operator and the best static subset from a frozen candidate set (“best-of-frozen-candidate-set”; the admissible universe $\binom{24}{6} = 134,596$ is *not* searched, so no claim about the best admissible static is possible). Interpretable only if G1-weak passes.

G2 (temporal order). The ordered schedule must differ from the median over frozen order permutations of the same epoch multiset (permutation-median comparator + cross-seed sign test). Interpretable only if G1-weak passes.

G3 (robustness). The weak advantage must persist under implemented stresses (heterogeneity, directedness, signed weights); PASS requires the faithful and mild-heterogeneity stages to pass.

G4 (identifiability). A past-only estimator must recover the active channel set from noisy observations of the surrogate (precision and recall > 0.6), track contraction (correlation > 0.5 , same-realization ground truth), and beat a factor-confounded baseline; tested under synchronous and asynchronous (LOCF) observation.

5 Synthetic Models and Experimental Design

FHN benchmark (G0A/G0B). Two ring layers of N FitzHugh–Nagumo oscillators ($\epsilon = 0.05$, $a_{\text{odd}} = 0.87$, $a_{\text{even}} = 0.97$, rotational coupling angle $\phi = \pi/2 - 0.1$, intra-layer coupling $\sigma_{\text{intra}} = 0.1$, nearest-neighbor ring) are coupled through a diffusive inter-layer term of strength σ_{12} on a switching subset of $N_{\text{IL}} = N/4$ mirror pairs, re-drawn uniformly at random every T_{swt} time units (RK4, $dt = 0.01$). Synchronization is declared when the tail time-average (final 25%) of the inter-layer error E_{12} falls below 0.02; both tolerances were frozen in the preregistration. G0B uses $N = 40$, $\sigma_{12} = 1.5$, total time 800, $T_{\text{swt}} \in \{2, \dots, 300\}$, five seeds per cell.

Minimal MSF system (G0C). A minimal $N = 2$ transverse-stability computation integrates the variational dynamics along the synchronized trajectory under smoothly switched anti-phase coupling channels ($\lambda_{\perp} = 2$, smoothing $\alpha = 5$), scanning $\sigma \in \{0.05, 0.15, 0.3, 0.5, 0.8\}$ and $T_{\text{swt}} \in \{11, 25, 100, 135\}$. The anti-phase channel prerequisite is verified before any science value is computed.

Linear stochastic surrogate (G1–G4). To probe mechanism transfer outside chaotic FHN dynamics we use a linear stochastic surrogate: two coupled layers of $N = 24$ nodes whose inter-layer difference evolves through a sequence of linear difference-step maps (local spectral radius 1.04, i.e. intentionally near-unstable between couplings; coupling gain $\kappa = 0.6$ on open channels; weak intra-layer mixing 0.06). The estimand is the contraction rate γ of the ordered product of the step maps over a frozen horizon $H = 240$: $\gamma > 0$ means the venue-difference (basis) process contracts. Schedules are *paired*: all comparators share the same base snapshots (six snapshots of six channel pairs) per seed, so differences isolate scheduling. The rate arms (fast / intermediate / slow) cycle the snapshot sequence 20/5/1 times within H . The switching arm and the best static candidate are selected on eight selection seeds and frozen *before* the eight disjoint evaluation seeds are opened; the static candidate set is the union of the base snapshots and a frozen 200-draw random search (248 candidates here).

Identifiability testbed (G4). On the same surrogate (horizon 600, fast dwell 6), a past-only, factor-free estimator based on venue differencing (window 8, observation noise 0.05) reconstructs the active channel set; it is compared against a deliberately factor-confounded level-correlation baseline (cf. [2, 13] for factor-removal rationale). Ground truth is the same realization’s schedule. Observation is either synchronous (1:1) or asynchronous with last-observation-carried-forward at ratio 1:3 — an *asynchrony/LOCF (Epps-like degradation)* design, not a causal demonstration of the Epps effect.

Custody. All grids, seeds, tolerances and decision rules above were frozen in preregistration v8 and its bound execution contract v6 before execution; the code and configuration were sealed by a content-addressed freeze manifest (v8), and the suite ran exactly once under an authorization token whose SHA-256 is recorded in every artifact (Appendix A).

Table 1: Preregistered gate outcomes of the single authorized cheap-suite run (attempt 1baa47da06fede2a, execution freeze v8). “Confirmatory” outcomes follow the frozen decision contract of Section 6; quantities marked “diagnostic” are reported but are, by the preregistered hierarchy, not gate outcomes. Zero technical failures, non-finite values or dropped seeds occurred in any gate.

Gate	Question (frozen)	Verdict	Status
G0A	exact paper-scale reproduction	NOT_RUN	outside cheap-suite scope
G0B	calibrated small- N demonstration	PASS	confirmatory (demonstration only)
G0C	minimal MSF T_{swt} -dependence	INCONCLUSIVE	confirmatory
G1-weak	switching beats equal-density static	INCONCLUSIVE (TIE)	confirmatory
G1-strict	beats average graph and best frozen static	NOT_INTERPRETABLE	diagnostic only
G2	temporal order beyond the multiset	NOT_INTERPRETABLE	diagnostic only
G3	robustness of the weak advantage	INCONCLUSIVE	confirmatory
G4	observational identifiability	FAIL	confirmatory

6 Statistical Decision Contract

Every paired comparison uses one shared, frozen rule. For paired differences $d_1, \dots, d_n$ ($n = 8$ evaluation seeds; differences with $|d| \leq 10^{-12}$ dropped and disclosed):

1. an exact two-sided sign test on the non-zero d_s at $\alpha = 0.05$ (with 8/8 equal signs, $p = 2 \cdot \binom{8}{0} / 2^8 = 0.0078$, so $n = 8$ can reach significance);
2. a preregistered *effect band* $b = \max(3 \hat{\sigma}_{\text{ddof}=1}, \text{floor})$, with frozen floors (0.02 for surrogate γ and FHN E_{12} ; 0.05 for the identifiability margin);
3. decision: if $p \geq \alpha \rightarrow$ INCONCLUSIVE (NOT_SIGNIFICANT); else if $\bar{d} > b \rightarrow$ PASS; else if $\bar{d} < -b \rightarrow$ FAIL; else $\rightarrow$ INCONCLUSIVE (TIE).

A bootstrap interval (2,000 resamples, frozen seed) is reported as a *descriptor only* and plays no role in decisions. A global failed-seed policy invalidates any cell with more than 20% technical failures; the executed run had **zero** failures, non-finite values or dropped seeds. The hierarchy is frozen: G1-strict and G2 are interpretable only if G1-weak passes; otherwise their computed values are reported as diagnostics and are not gate outcomes.

7 Results

Table 1 and Figure 1 summarize the outcomes of the single authorized run (sealed attempt 1baa47da06fede2a); Table 3 lists every decision statistic.

G0B: PASS (confirmatory; calibrated demonstration only). All fast cells ($T_{\text{swt}} \leq 10$) synchronize (fraction 1.0) and all slow cells ($T_{\text{swt}} \geq 160$) do not (fraction 0.2), with a monotone descent in between (Table 2, Figure 2). Five of five seeds succeeded in every cell. This demonstrates the switching-synchronization mechanism in a calibrated small- N setting; it is *not* an exact reproduction of the paper regime and supports no claim about the N -dependence of the switching threshold.

G0C: INCONCLUSIVE (confirmatory). Transverse contraction exists: $\Psi < 0$ across the entire frozen grid (at the fastest $T_{\text{swt}} = 11$: $\Psi = -0.008$ at $\sigma = 0.05$ down to -0.177 at $\sigma = 0.8$). But the stability onset is $\sigma = 0.05$ — the smallest grid value — at *every* $T_{\text{swt}} \in \{11, 25, 100, 135\}$, so no

Preregistered gate outcomes (attempt 1baa47da06fede2a, freeze v8)

G0A exact reproduction	NOT_RUN	outside cheap-suite scope
G0B calibrated demonstration	PASS	calibrated demonstration only
G0C minimal MSF	INCONCLUSIVE	no T_{swt} dependence resolved
G1-weak switching vs static	INCONCLUSIVE (TIE)	8/8 positive, below band
G1-strict vs avg/best static	NOT_INTERPRETABLE	diagnostic: 8/8 negative
G2 temporal order	NOT_INTERPRETABLE	diagnostic: not significant
G3 robustness	INCONCLUSIVE	positive but inside bands
G4 identifiability	FAIL	P/R $\ll$ 0.6; asynchrony collapse

Figure 1: Gate summary. Grey chips mark quantities that the frozen hierarchy declares non-interpretable (diagnostic only).

Table 2: G0B calibrated demonstration ($\sigma_{12} = 1.5$, $N = 40$, $N_{\text{IL}} = 10$, 5 seeds per cell, zero failures). The frozen criterion constrains only the fast band ($T_{\text{swt}} \leq 10$: all fractions ≥ 0.5) and the slow band ($T_{\text{swt}} \geq 160$: all fractions < 0.5); intermediate cells are reported descriptively.

T_{swt}	2	5	10	20	40	80	160	300
fraction synchronized	1.0	1.0	1.0	1.0	1.0	0.6	0.2	0.2

switching-time dependence was resolved on the frozen grid (PARTIAL_NO_TSWT_DEPENDENCE). Descriptively, contraction is present; confirmatorily, the T_{swt} -dependence that the translation needs was not shown at this resolution.

G1-weak: INCONCLUSIVE / TIE (confirmatory). The rate arm selected on the frozen selection seeds was *fast* (selection scores: fast +0.0366, intermediate +0.0071, slow -0.0155). On the eight disjoint evaluation seeds, all eight paired differences $\gamma_{\text{arm}} - \gamma_{\text{static}}$ are positive (sign test $p = 0.0078$), but the mean +0.0420 lies inside the preregistered band $b = 0.0666$ (Figure 3, left). Under the frozen rule this is INCONCLUSIVE (TIE): **directional evidence exists, but the effect failed the preregistered magnitude criterion.** We do not reinterpret this as a pass.

G1-strict and G2: NOT_INTERPRETABLE (diagnostic only). Because G1-weak did not pass, the frozen hierarchy renders G1-strict and G2 non-interpretable; we report their values strictly as diagnostics (Figure 3, middle and right; Table 3). Diagnostically, the G1-strict differences $\gamma_{\text{arm}} - \max(\gamma_{\text{avg}}, \gamma_{\text{best static}})$ are negative in *all eight* evaluation seeds (mean -0.0046): the switching arm never beat the aggregate comparator in this ensemble. The G2 order statistic is small (+0.0030) and not significant ($p = 0.070$). These diagnostics must not be promoted to confirmatory conclusions — in either direction.

G3: INCONCLUSIVE (confirmatory). Every stage shows a positive mean advantage $\gamma_{\text{fast}} - \gamma_{\text{static}}$ (from +0.0455 to +0.0631), with 8/8 positive signs in four of five stages — yet no stage exceeds its frozen band (Figure 4, Table 3), and the gate requires the faithful and mild-heterogeneity




Figure 2: G0B synchronization fraction versus switching period (confirmatory). The frozen criterion constrains only the shaded bands.

Table 3: Paired decision statistics ($n = 8$ evaluation seeds per test; exact two-sided sign test on non-zero differences; effect band $b = \max(3 \hat{\sigma}_{\text{ddof}=1}, \text{floor})$; floors: 0.02 for surrogate γ , 0.05 for the identifiability margin). Rows marked $\dagger$ are diagnostic (NOT INTERPRETABLE under the frozen hierarchy).

Test	mean $\bar{d}$	band b	signs (+/-)	p	verdict
G1-weak: $\gamma_{\text{arm}} - \gamma_{\text{static}}$	+0.0420	0.0666	8/0	0.0078	INCONCLUSIVE (TIE)
G1-strict $\dagger$: $\gamma_{\text{arm}} - \max(\gamma_{\text{avg}}, \gamma_{\text{best}})$	-0.0046	0.0200	0/8	0.0078	(diagnostic)
G2 $\dagger$: $\gamma_{\text{ord}} - \text{med}(\gamma_{\text{perm}})$	+0.0030	0.0200	7/1	0.0703	(diagnostic)
G3 faithful	+0.0492	0.0789	8/0	0.0078	INCONCLUSIVE (TIE)
G3 mild heterogeneity	+0.0455	0.1715	8/0	0.0078	INCONCLUSIVE (TIE)
G3 strong heterogeneity	+0.0631	0.2984	6/2	0.2891	INCONCLUSIVE (NOT.SIG.)
G3 directed	+0.0603	0.1006	8/0	0.0078	INCONCLUSIVE (TIE)
G3 signed	+0.0467	0.0799	8/0	0.0078	INCONCLUSIVE (TIE)
G4 beats-baseline margin	+0.1625	0.0646	8/0	0.0078	PASS (condition; gate FAILs on P/R)

stages to pass. The signed stage satisfied its implementation prerequisites (per-seed negative weights present; Frobenius budget matched to 10^{-9}), so the verdict is a clean INCONCLUSIVE, not an execution artifact.

G4: FAIL (confirmatory). The factor-free basis estimator recovers the active channel set with precision = recall = 0.353 under synchronous observation and 0.254 under asynchronous/LOCF observation — far below the preregistered 0.6 bar (Figure 5). It does beat the factor-confounded baseline (margin +0.162, band 0.065, 8/8 positive, $p = 0.0078$) and, under synchronous sampling, tracks contraction well (correlation 0.948). Under asynchronous/LOCF sampling the contraction correlation collapses to 0.010 — an Epps-like degradation [10, 29], consistent with the classical observation that asynchronous sampling destroys short-horizon comovement estimates even when consistent estimators exist [19]. Two of the four frozen conditions fail, so the gate fails.




Figure 3: Per-seed paired differences with preregistered effect bands (red). Left: confirmatory G1-weak. Middle/right (grey background): diagnostic G1-strict and G2, NOT_INTERPRETABLE under the frozen hierarchy.

G0A: NOT_RUN (permanent decision). The exact paper-scale reproduction ($\sigma_{12} = 0.1$, N up to 400, 4,000 time units per cell) is outside the authorized cheap-suite scope and is recorded permanently as NOT_RUN — *outside the cheap-suite scope and unnecessary for the financial-translation decision*. It could be conducted only as a separate replication project; whatever its outcome, it cannot rescue G1–G4 (which concern translation, order effects and identifiability, not the mechanism’s existence) and cannot authorize real-market research under the closed preregistration.

8 Why Aggregate Connectivity Dominates the Interpretation

The pattern of outcomes is exactly what fast-switching averaging theory predicts for this ensemble. The selected arm was the *fastest*; blinking and fast-switching results [4, 18, 30] say that in this regime the switched system inherits the stability of the time-averaged network, with finite-rate corrections. Consistently: the weak advantage over an equal-density *static sparse* graph is directionally positive (a static sparse subset is simply a worse average), while the diagnostic strict comparison against the average-graph operator is negative in every seed. Nothing in the confirmatory record shows switching delivering anything *beyond* its time-average here.

Two disciplined qualifications. First, this is not an impossibility claim: designed temporal networks can provably outperform static ones under resource constraints [31], and our frozen ensemble was not designed to sit in that regime; the negative reading applies to this preregistered ensemble, not to all switching systems. Second, the G2 permutation-median comparator aggregates signed order effects across seeds and can cancel sign-varying order effects; this limitation was frozen and disclosed in advance.

For finance the implication is direct: under this translation, measured “switching structure” carried no confirmed information beyond the average network. Aggregate connectedness measures [5, 7, 8] — and time-varying correlation models [9] with the usual heteroskedasticity discipline [14] — already target that average object. A switching-based market study would therefore need to demonstrate, first and preregistered, an effect beyond the average graph; ours did not.




Figure 4: G3 per-stage paired differences with per-stage effect bands (confirmatory; all stages INCONCLUSIVE).

9 Observational Identifiability and Asynchrony

Even in the most favorable conditions finance can never offer — known synthetic dynamics, same-realization ground truth, a factor-free estimator by construction, and synchronous noise-controlled observations — channel-set recovery reached precision and recall of only 0.353. Beating the factor-confounded baseline (margin +0.162; the baseline’s precision is 0.190) confirms that venue differencing removes the common factor, in the spirit of factor-removal approaches [2, 13], but “better than a confounded baseline” is far from “reliable”. Under asynchronous/LOCF observation at a modest 1:3 ratio, recovery degrades further (0.254) and the contraction correlation collapses from 0.948 to 0.010. We stress the frozen terminology: this is an *asynchrony/LOCF* (*Epps-like degradation*) experiment — it reproduces the practical consequence highlighted by Epps [10] in our setting; it is not a causal demonstration of the Epps effect, and estimators designed for non-synchronous data [19, 29] were not the object under test. Since real markets add unknown dynamics, latent factors, microstructure noise and far harsher asynchrony, the preregistered conclusion is that the mechanism’s observational signature is not reliably identifiable in this class of designs; a real-data “detection” built on such an estimator would not be interpretable.

10 Limitations

1. **Scope of the negative result.** All conclusions are relative to the frozen ensemble, grids and decision contract. We did not search parameter space for regimes where switching beats the average [31]; that would be a different, new preregistration.
2. **Small n , wide bands.** With $n = 8$ paired seeds and a $3\hat{\sigma}$ band, magnitude criteria are conservative; consistent directional effects (8/8 signs) can land INCONCLUSIVE, as G1-weak and four G3 stages did. This is by design (anti-inflation), but it limits power.
3. **Comparator strength.** G1-strict uses the best of a frozen 248-candidate static set, not the best admissible static subset ($\binom{24}{6} = 134,596$); claims about the true static optimum are impossible and none are made.




Figure 5: G4 identifiability metrics under synchronous (1:1) and asynchronous/LOCF (1:3) observation (confirmatory; gate FAIL).

4. **G2 statistic.** The permutation-median comparator detects a systematic displacement from the permutation median; sign-varying order effects across seeds can cancel (frozen, disclosed limitation).
5. **G0C resolution.** The onset grid starts at $\sigma = 0.05$ and the onset saturated at that smallest value for every T_{swt} ; a finer grid might resolve dependence. The minimal $N = 2$ construction is also the weakest link to the full spatiotemporal system.
6. **G0A not run.** The exact paper-scale reproduction was not executed; G0B is a calibrated demonstration at different coupling and size and cannot substitute for it.
7. **The surrogate is not a market model.** It is a mechanism testbed. No claim about real markets, trading, price prediction or market causality follows from any result here — positive or negative. No real market data were used anywhere.

11 Reproducibility and Audit Trail

The study enforced preregistration-grade custody end to end: (i) scientific preregistration and operational execution contract, versioned and bound by canonical hashes (v8/v6 at execution); (ii) a content-addressed freeze manifest over all code, configuration, tests and documents (117 files), whose stored content hash equals independent recomputation; (iii) a single authorized execution from the freeze commit with a clean tree, identified by an authorization token recorded only as its SHA-256, producing attempt `1baa47da06fede2a`; (iv) a sealed attempt directory whose manifest inventories every artifact by size and SHA-256 and whose seal equals the manifest content hash; (v) an independent audit that reconstructed the package and recomputed all identities, hashes and decision statistics exactly; (vi) byte-for-byte custody of the sealed outputs in the repository under an annotated tag, distinct from the execution freeze commit. Every figure and table in this paper is generated from the frozen JSON reports without re-running any simulation. Earlier execution freezes v2–v7 were each found non-executable by independent audits and were *never run*; consequently no result predates the sealed attempt. All identifiers appear in Appendix A.

12 Conclusion

SwitchSync Markets demonstrates the switching-synchronization mechanism in a calibrated synthetic setting (G0B PASS), but under the preregistered criteria it does not establish an advantage beyond aggregate connectivity (G1-weak INCONCLUSIVE/TIE; G1-strict diagnostic and non-interpretable), a robust temporal-order effect (G2 non-interpretable; G3 INCONCLUSIVE), or reliable observational identifiability (G4 FAIL). The financial translation is therefore not supported for empirical or operational development under the preregistered criteria: no predictive or operational market study is justified, and no trading signal, backtest, PnL, Sharpe or deployment claim exists or is implied. No real market data were used. This outcome does not refute the mathematical results of the seed paper [12]; it delimits what that mechanism can currently be expected to deliver as a measurable, decision-relevant object in financial data, and it documents — with full custody — how a preregistered negative result of this kind can be produced and audited.

A Audit and Reproducibility Appendix

Identifiers. Table 4 lists the immutable hashes. The execution scope was `cheap-suite`; G0A is permanently NOT-RUN. The authorization token is recorded only as its SHA-256:

3b4b903518fa5cb337cbf5416aa8cc6db8db47c1548b39545f562158b07aca29

the attempt id satisfies `attempt = sha256(campaign || scope || token_sha)[: 16]`, so an auditor recomputes it without the raw token. Zero failure records, non-finite values or dropped seeds occurred. No real data were accessed.

Table 4: Immutable identifiers of the preregistration, the execution freeze and the sealed result. Earlier execution freezes v2–v7 were each found non-executable by independent audit and were *never run*; no result predates the attempt below.

Object	SHA-256 / identifier
prereg v8 (canonical)	a667434d0d7c5905d62cf81b79699f2799e2597128631cab1e4792bed6053dd9
prereg v8 (file)	896163891b8b43f42511a1a57294d6075f3b65f33877780daf82cc03a54d568e
execution contract v6 (canonical)	6dbe79ddf1d0b6e5e63f74a5ae3ed947abfc837419f4b7938fdcf8c09f525ee8
execution contract v6 (file)	80609ef5a90817813f0a9978ec017ab00734c3681abb0429fbe4ec4582b8c2bb
execution freeze v8 (content hash)	2991e064ced6154c149fa57d4fc77e855bede51750565c264972b369778d112d
freeze commit (= runtime HEAD)	5ca40c7da87f583c7cea1cefd00cc978c3731d98
result-custody commit	48339e47d3a0911949cf25c83a50e63fc9572aa5
campaign / attempt id	44be3a5d07e67b5e / 1baa47da06fede2a
manifest content hash (= SEALED)	4e029d814437695d99c4b5cc62cbe3cf94ebd0e46c6e9ae51d54cc5f28f3c33d
g0b_calibrated_v2.json	d592b75d59b6e50c51ec6263c0a8469654651e0ef247fd2fadd3cf5adad4d4e5
g0c_msf_v2.json	150a53dc42236742d347012704a00a4d2e1db52d99bb800341bbc3ec6de484d0
g1_g2_paired_v2.json	b6242cd72510ce6d8b0daa5edba31b1c124f3fa519eae92ac0a408f06decbce0
g3_robustness_v2.json	2f16a4578c5b48cbb8621455b5a37042ec615a7127c9d709acb43ab1a3414964
g4_identifiability_v2.json	bd148d1ed23ba057df41dedfbd7414b7d450a7f5c28b59617e354d5bbb597e6

Seed blocks (frozen in prereg v8). chaos {7}; G0A {11, 12, 13}; G0B {11, ..., 15}; G0C {31}; paired selection {41, ..., 48}; paired evaluation {81, ..., 88}; G3 stages {51, ..., 58}; identifiability {61, ..., 68}.

Environment (pinned and verified at execution). Python 3.14.3; numpy 2.4.4; scipy 1.17.1; matplotlib 3.10.9; Linux x86-64 (glibc 2.35); RK4 with $dt = 0.01$; runtime ≈ 8 min 44 s for the full cheap suite.

Execution history declaration. Execution freezes v2–v7 were each found non-executable by independent pre-execution audits (defects ranged from custody flaws to a missing frozen tolerance block) and were superseded without ever producing a result. The sealed attempt `1baa47da06fede2a` under freeze v8 is the *only* scientific execution of this project.

B Frozen Decision Details

Shared rule. Exact two-sided sign test on non-zero paired differences at $\alpha = 0.05$; effect band $b = \max(3\hat{\sigma}_{\text{ddof}=1}, \text{floor})$; PASS/FAIL only outside $\pm b$; bootstrap (2,000 resamples, frozen seed 20260713) reported as descriptor only.

G1/G2 selection protocol. Arm and best-static comparator selected on the selection seed block only, using a common per-seed validity mask over *all* rate arms and *all* 248 static candidates, then frozen before evaluation seeds are opened. Ties break canonically (arm order fast < intermediate < slow; lexicographic subsets).

G4 conditions (all four required). mean precision > 0.6 ; mean recall > 0.6 ; same-realization contraction correlation > 0.5 ; significant positive margin over the factor-confounded baseline (floor 0.05). Observed: 0.353; 0.353; 0.948 (synchronous); margin +0.162 (PASS) — two conditions fail, hence gate FAIL.

References

- [1] C. Alexander, J. Choi, H. Park, and S. Sohn. Bitmex bitcoin derivatives: Price discovery, informational efficiency, and hedging effectiveness. *Journal of Futures Markets*, 40:23, 2020. doi: 10.1002/fut.22050.
- [2] M. Barigozzi and M. Hallin. A network analysis of the volatility of high-dimensional financial series. *Journal of the Royal Statistical Society: Series C*, 66:581, 2017. doi: 10.1111/rssc.12177.
- [3] D. G. Baur and T. Dimpfl. Price discovery in bitcoin spot or futures? *Journal of Futures Markets*, 39:803, 2019. doi: 10.1002/fut.22004.
- [4] I. V. Belykh, V. N. Belykh, and M. Hasler. Blinking model and synchronization in small-world networks with a time-varying coupling. *Physica D*, 195:188, 2004. doi: 10.1016/j.physd.2004.03.013.
- [5] M. Billio, M. Getmansky, A. W. Lo, and L. Pelizzon. Econometric measures of connectedness and systemic risk in the finance and insurance sectors. *Journal of Financial Economics*, 104: 535, 2012. doi: 10.1016/j.jfineco.2011.12.010.
- [6] S. Boccaletti, D.-U. Hwang, M. Chavez, A. Amann, J. Kurths, and L. M. Pecora. Synchronization in dynamical networks: Evolution along commutative graphs. *Physical Review E*, 74: 016102, 2006. doi: 10.1103/PhysRevE.74.016102.

- [7] F. X. Diebold and K. Yilmaz. Measuring financial asset return and volatility spillovers, with application to global equity markets. *The Economic Journal*, 119:158, 2009. doi: 10.1111/j.1468-0297.2008.02208.x.
- [8] F. X. Diebold and K. Yilmaz. On the network topology of variance decompositions: Measuring the connectedness of financial firms. *Journal of Econometrics*, 182:119, 2014. doi: 10.1016/j.jeconom.2014.04.012.
- [9] R. Engle. Dynamic conditional correlation. *Journal of Business & Economic Statistics*, 20:339, 2002. doi: 10.1198/073500102288618487.
- [10] T. W. Epps. Comovements in stock prices in the very short run. *Journal of the American Statistical Association*, 74:291, 1979. doi: 10.1080/01621459.1979.10482508.
- [11] M. C. Eser, E. S. Medeiros, M. Riza, and A. Zakharova. Edges of inter-layer synchronization in multilayer networks with time-switching links. *Chaos*, 31:103119, 2021. doi: 10.1063/5.0065310.
- [12] M. C. Eser, E. S. Medeiros, M. Riza, and M. Engel. Dynamic link switching induces stable synchronized states in sparse networks, 2025. URL <https://arxiv.org/abs/2507.08007>. arXiv:2507.08007 [nlin.AO].
- [13] J. Fan, Y. Liao, and M. Mincheva. Large covariance estimation by thresholding principal orthogonal complements. *Journal of the Royal Statistical Society: Series B*, 75:603, 2013. doi: 10.1111/rssb.12016.
- [14] K. J. Forbes and R. Rigobon. No contagion, only interdependence: Measuring stock market comovements. *The Journal of Finance*, 57:2223, 2002. doi: 10.1111/0022-1082.00494.
- [15] D. Ghosh, M. Frasca, A. Rizzo, S. Majhi, S. Rakshit, K. Alfaro-Bittner, and S. Boccaletti. The synchronized dynamics of time-varying networks. *Physics Reports*, 949:1, 2022. doi: 10.1016/j.physrep.2021.10.006.
- [16] J. Gonzalo and C. W. J. Granger. Estimation of common long-memory components in cointegrated systems. *Journal of Business & Economic Statistics*, 13:27, 1995. doi: 10.1080/07350015.1995.10524576.
- [17] J. Hasbrouck. One security, many markets: Determining the contributions to price discovery. *The Journal of Finance*, 50:1175, 1995. doi: 10.1111/j.1540-6261.1995.tb04054.x.
- [18] M. Hasler, V. Belykh, and I. Belykh. Dynamics of stochastically blinking systems. part i: Finite time properties. *SIAM Journal on Applied Dynamical Systems*, 12:1007, 2013. doi: 10.1137/120893409.
- [19] T. Hayashi and N. Yoshida. On covariance estimation of non-synchronously observed diffusion processes. *Bernoulli*, 11:359, 2005. doi: 10.3150/bj/1116340299.
- [20] P. Holme and J. Saramäki. Temporal networks. *Physics Reports*, 519:97, 2012. doi: 10.1016/j.physrep.2012.03.001.
- [21] A. Jadbabaie, J. Lin, and A. S. Morse. Coordination of groups of mobile autonomous agents using nearest neighbor rules. *IEEE Transactions on Automatic Control*, 48:988, 2003. doi: 10.1109/TAC.2003.812781.

- [22] H. Jardon-Kojakhmetov, C. Kuehn, and I. P. Longo. Persistent synchronization of heterogeneous networks with time-dependent linear diffusive coupling. *SIAM Journal on Applied Dynamical Systems*, 23:1540, 2024. doi: 10.1137/23M1602024.
- [23] V. Kohar, P. Ji, A. Choudhary, S. Sinha, and J. Kurths. Synchronization in time-varying networks. *Physical Review E*, 90:022812, 2014. doi: 10.1103/PhysRevE.90.022812.
- [24] L. Moreau. Stability of multiagent systems with time-dependent communication links. *IEEE Transactions on Automatic Control*, 50:169, 2005. doi: 10.1109/TAC.2004.841888.
- [25] R. Olfati-Saber and R. M. Murray. Consensus problems in networks of agents with switching topology and time-delays. *IEEE Transactions on Automatic Control*, 49:1520, 2004. doi: 10.1109/TAC.2004.834113.
- [26] L. M. Pecora and T. L. Carroll. Synchronization in chaotic systems. *Physical Review Letters*, 64:821, 1990. doi: 10.1103/PhysRevLett.64.821.
- [27] L. M. Pecora and T. L. Carroll. Master stability functions for synchronized coupled systems. *Physical Review Letters*, 80:2109, 1998. doi: 10.1103/PhysRevLett.80.2109.
- [28] W. Ren and R. W. Beard. Consensus seeking in multiagent systems under dynamically changing interaction topologies. *IEEE Transactions on Automatic Control*, 50:655, 2005. doi: 10.1109/TAC.2005.846556.
- [29] M. Scholes and J. Williams. Estimating betas from nonsynchronous data. *Journal of Financial Economics*, 5:309, 1977. doi: 10.1016/0304-405X(77)90041-1.
- [30] D. J. Stilwell, E. M. Bollt, and D. G. Roberson. Sufficient conditions for fast switching synchronization in time-varying network topologies. *SIAM Journal on Applied Dynamical Systems*, 5:140, 2006. doi: 10.1137/050625229.
- [31] Y. Zhang and S. H. Strogatz. Designing temporal networks that synchronize under resource constraints. *Nature Communications*, 12:3273, 2021. doi: 10.1038/s41467-021-23446-9.
- [32] J. Zhou, X. Guo, Z. Liu, Y.-C. Lai, and W. Lin. Random temporal connections promote network synchronization. *Physical Review E*, 100:032302, 2019. doi: 10.1103/PhysRevE.100.032302.